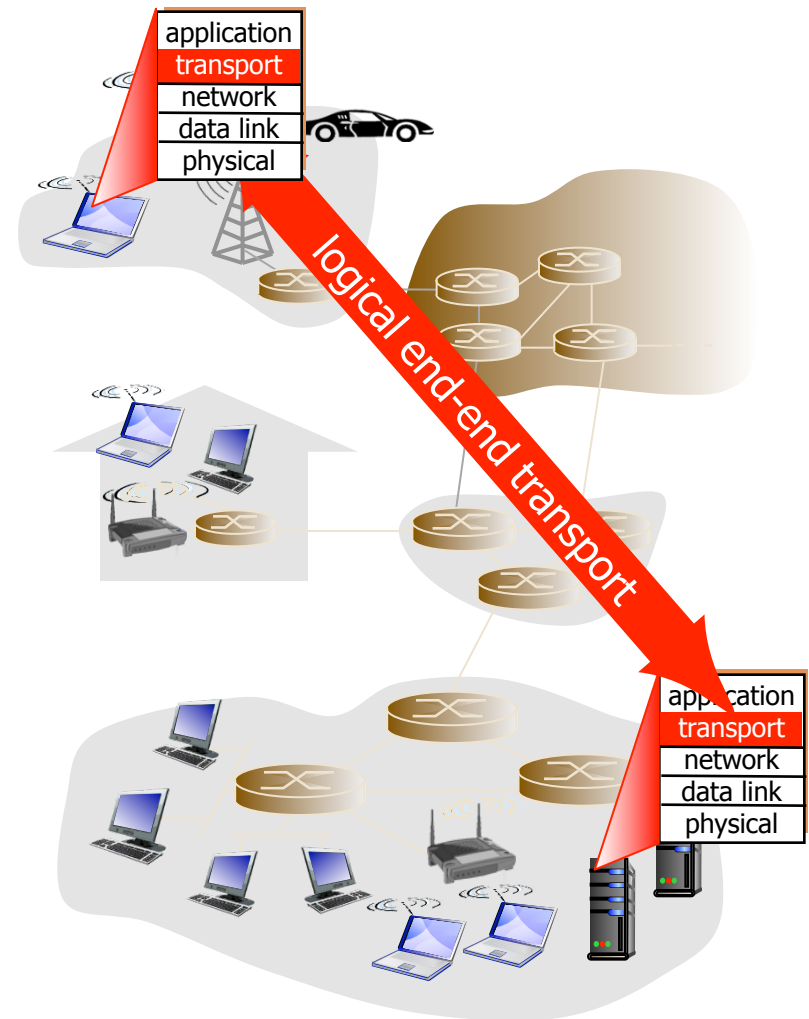


TRANSPORT LAYER

Connection Time

Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
 - ▣ send side: breaks app messages into *segments*, passes to network layer
 - ▣ rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
 - ▣ Internet: TCP and UDP



Transport vs. network layer

- *network layer*: logical communication between hosts
- *transport layer*: logical communication between processes
 - ▣ relies on, enhances, network layer services

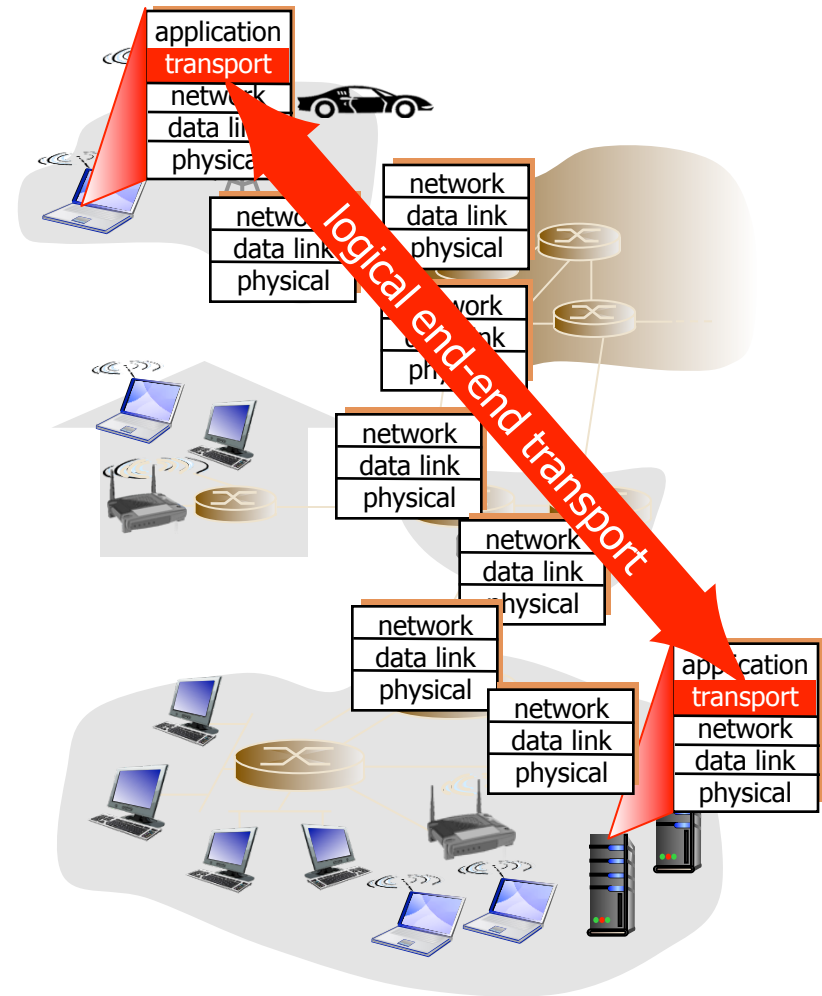
household analogy:

5 kids in Ann's house sending letters to 5 kids in Bill's house:

- hosts = houses
- processes = kids
- app messages = letters in envelopes
- transport protocol = Ann and Bill who demux to in-house siblings
- network-layer protocol = postal service

Internet transport-layer protocols

- reliable, in-order delivery (TCP)
 - ▣ congestion control
 - ▣ flow control
 - ▣ connection setup
- unreliable, unordered delivery: UDP
 - ▣ no-frills extension of “best-effort” IP
- services not available:
 - ▣ delay guarantees
 - ▣ bandwidth guarantees



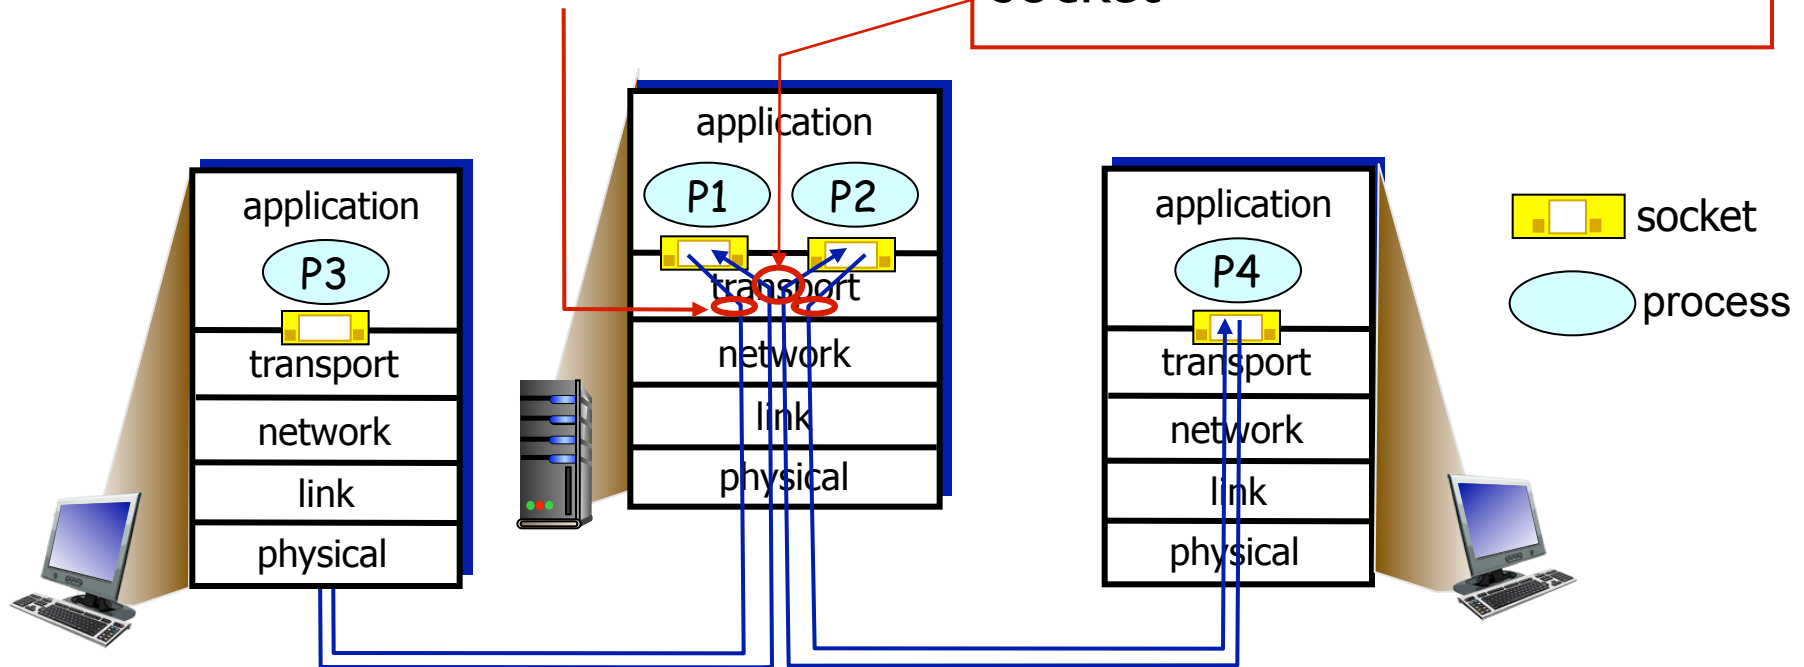
Multiplexing/demultiplexing

multiplexing at sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)

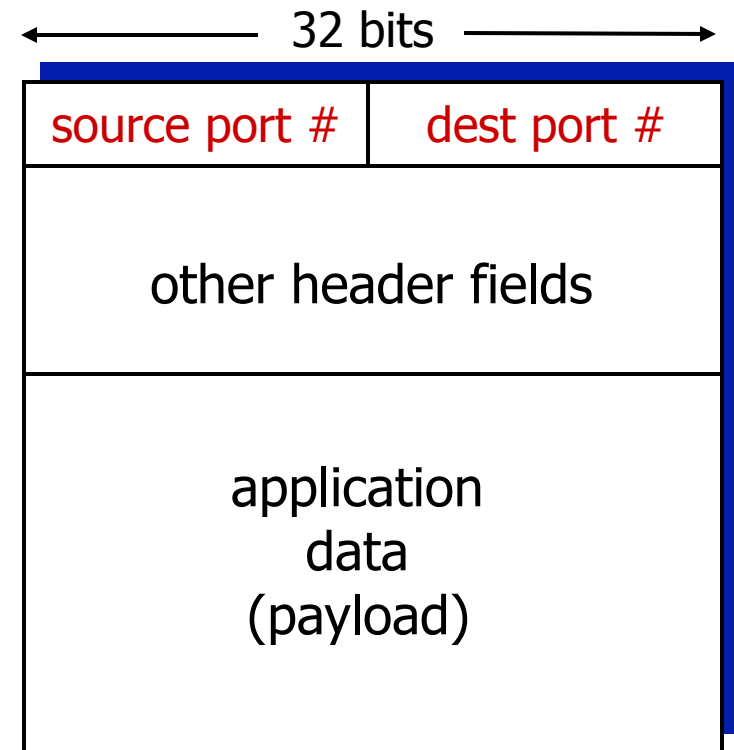
demultiplexing at receiver:

use header info to deliver received segments to correct socket



How demultiplexing works

- host receives IP datagrams
 - ▣ each datagram has source IP address, destination IP address
 - ▣ each datagram carries one transport-layer segment
 - ▣ each segment has source, destination port number
- host uses *IP addresses & port numbers* to direct segment to appropriate socket



TCP/UDP segment format

Connectionless demultiplexing

recall: when creating datagram to send into UDP socket, must specify

- destination IP address
- destination port #

□ when host receives UDP segment:

- ▣ checks destination port # in segment
- ▣ directs UDP segment to socket with that port #

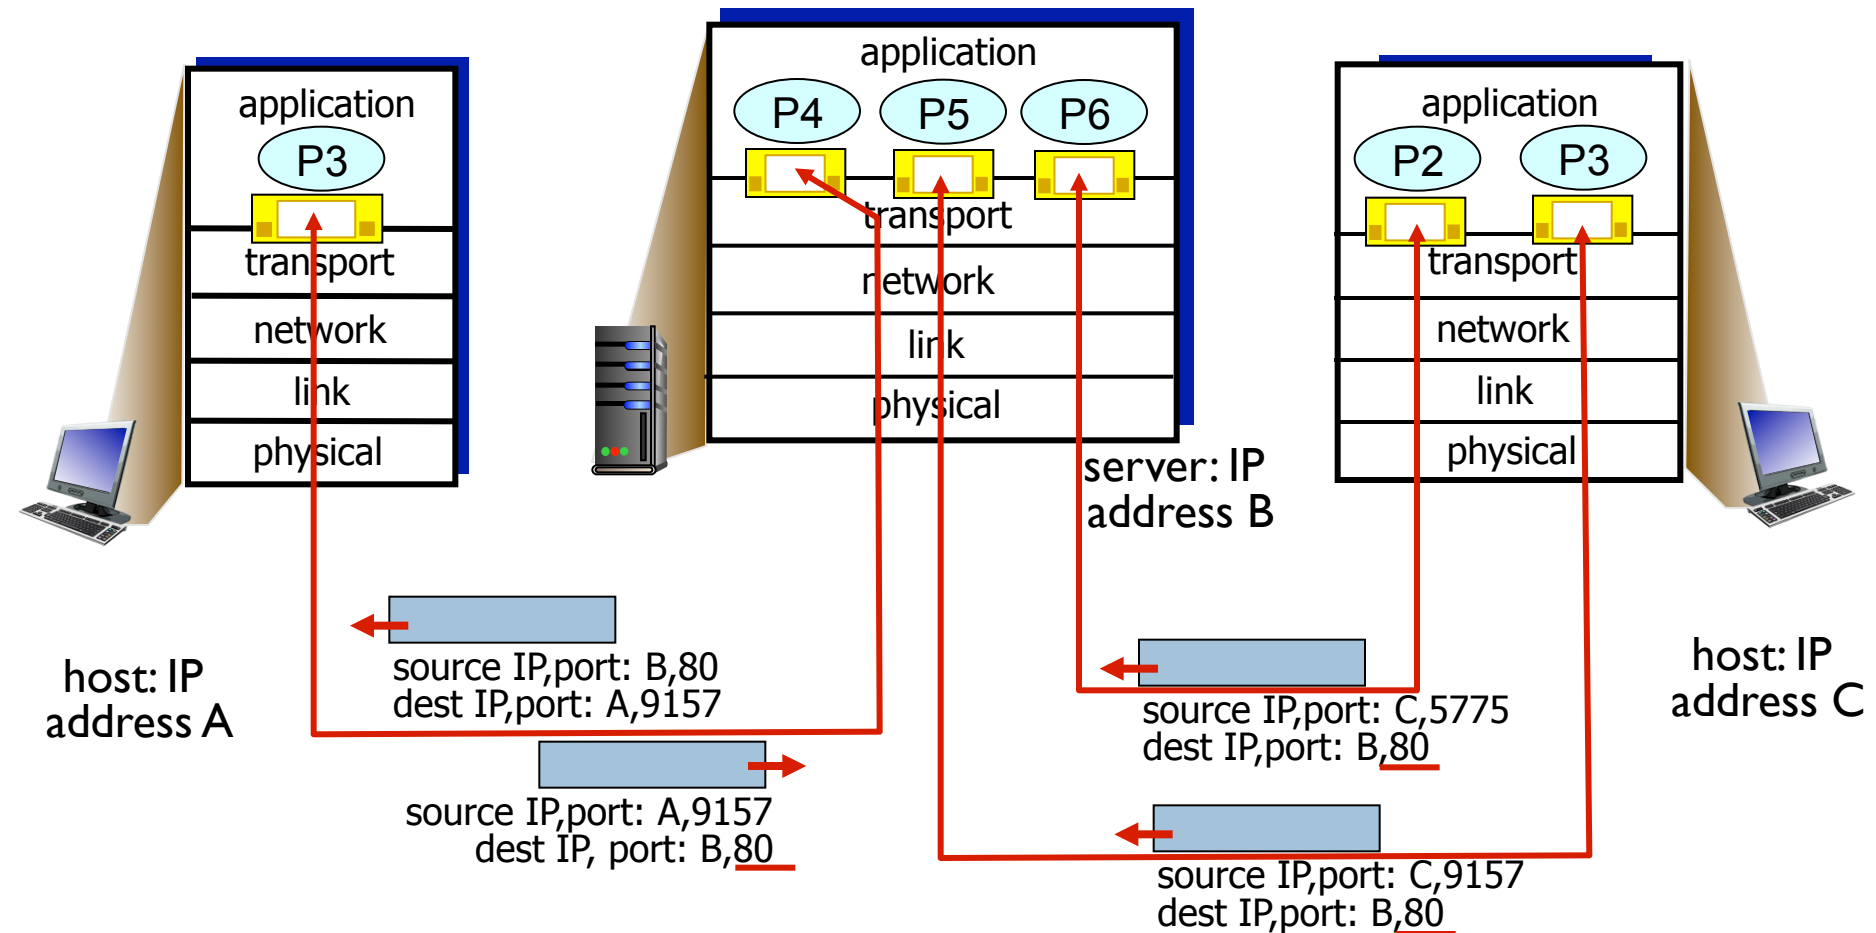


IP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at dest

Connection-oriented demux

- TCP socket identified by 4-tuple:
 - ▣ source IP address
 - ▣ source port number
 - ▣ dest IP address
 - ▣ dest port number
- demux: receiver uses all four values to direct segment to appropriate socket
- server host may support many simultaneous TCP sockets:
 - ▣ each socket identified by its own 4-tuple
- web servers have different sockets for each connecting client
 - ▣ non-persistent HTTP will have different socket for each request

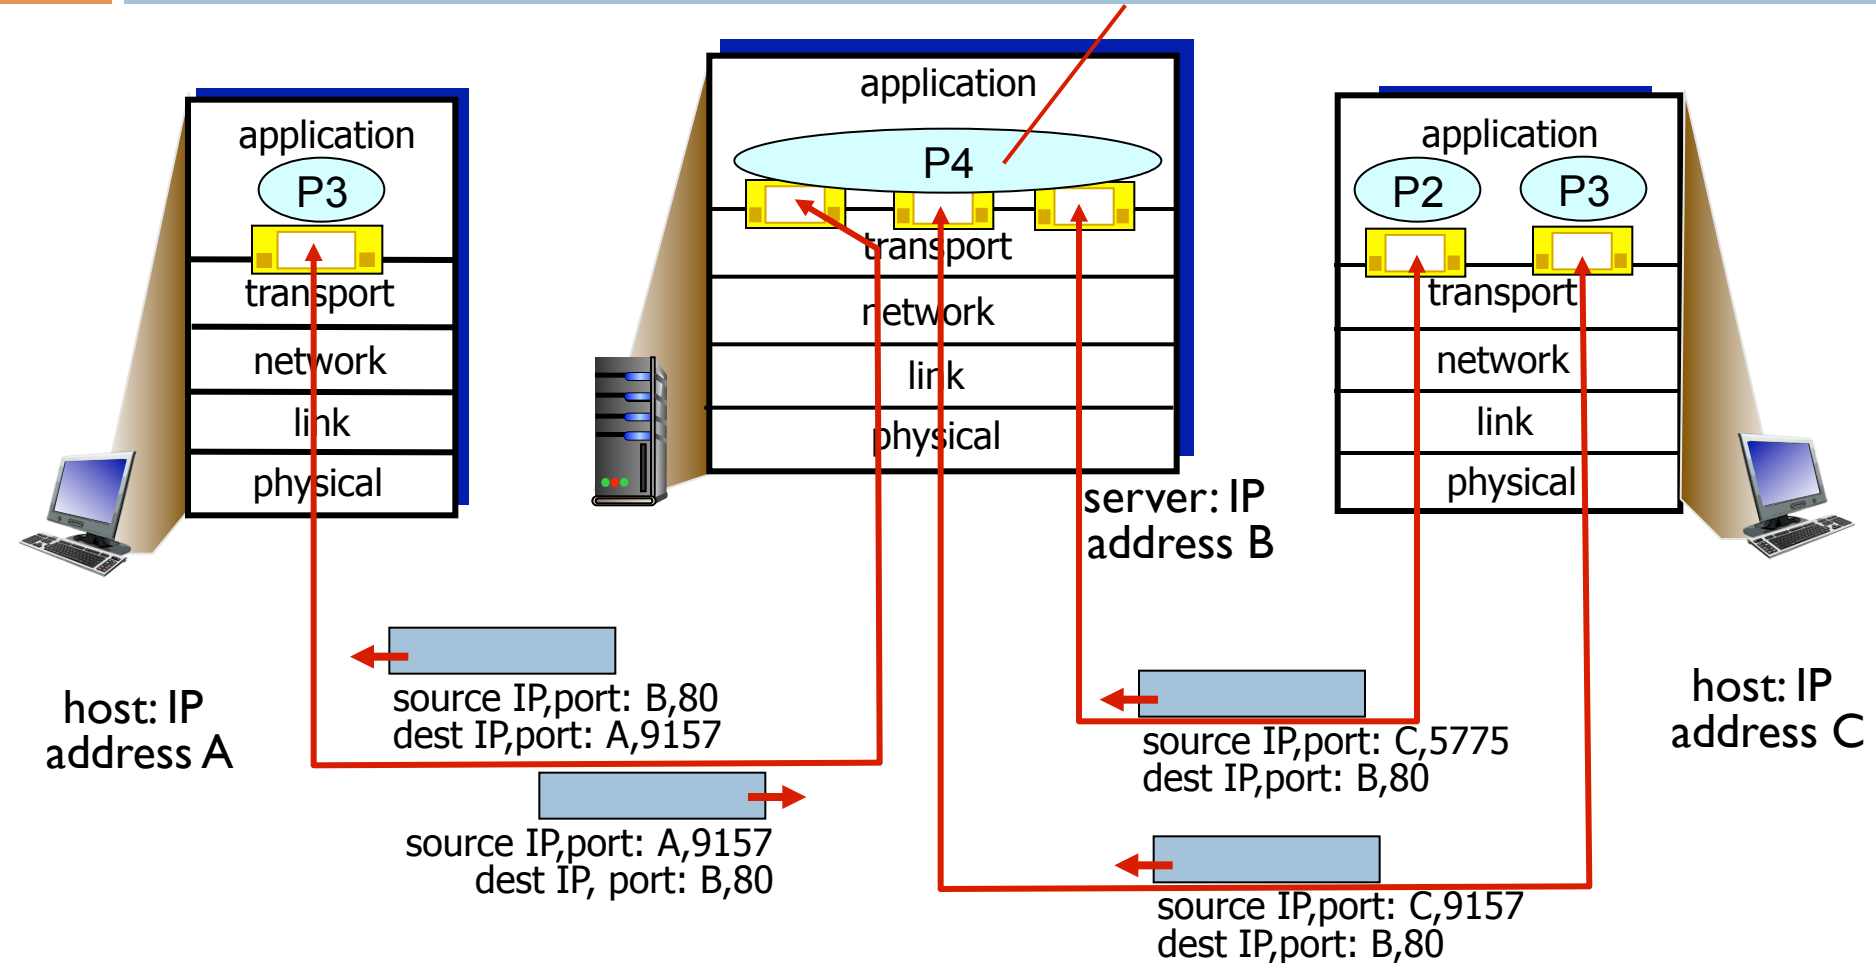
Connection-oriented demux: example



three segments, all destined to IP address: B,
dest port: 80 are demultiplexed to *different* sockets

Connection-oriented demux: example

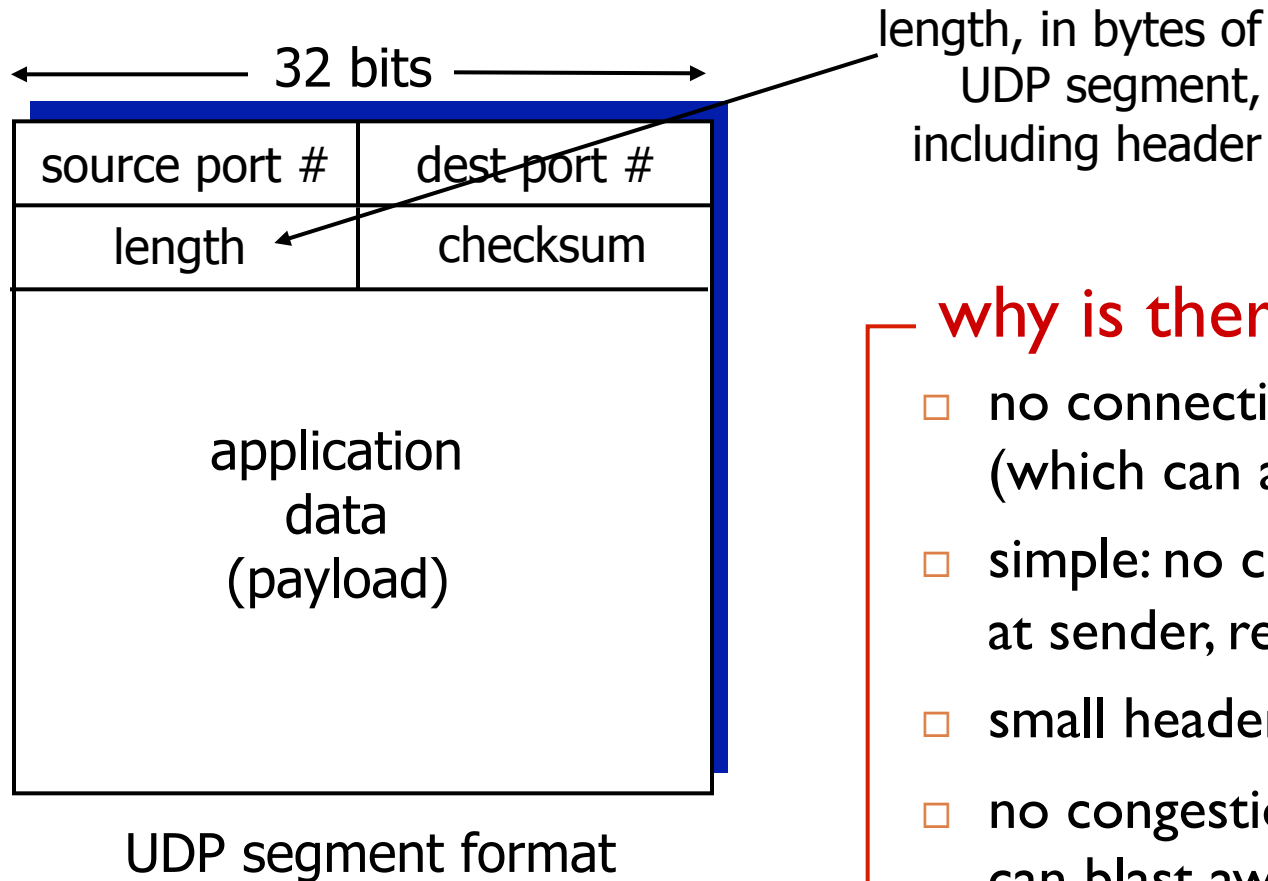
threaded server



UDP: User Datagram Protocol [RFC 768]

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
 - ▣ lost
 - ▣ delivered out-of-order to app
- *connectionless*:
 - ▣ no handshaking between UDP sender, receiver
 - ▣ each UDP segment handled independently of others
- ▣ UDP use:
 - ▣ streaming multimedia apps (loss tolerant, rate sensitive)
 - ▣ DNS
 - ▣ SNMP
- ▣ reliable transfer over UDP:
 - ▣ add reliability at application layer
 - ▣ application-specific error recovery!

UDP: segment header



why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control: UDP can blast away as fast as desired

UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

sender:

- treat segment contents, including header fields, as sequence of 16-bit integers
- checksum: addition (one's complement sum) of segment contents
- sender puts checksum value into UDP checksum field

receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - ▣ NO - error detected
 - ▣ YES - no error detected. *But maybe errors nonetheless?*
More later

Internet checksum: example

example: add two 16-bit integers

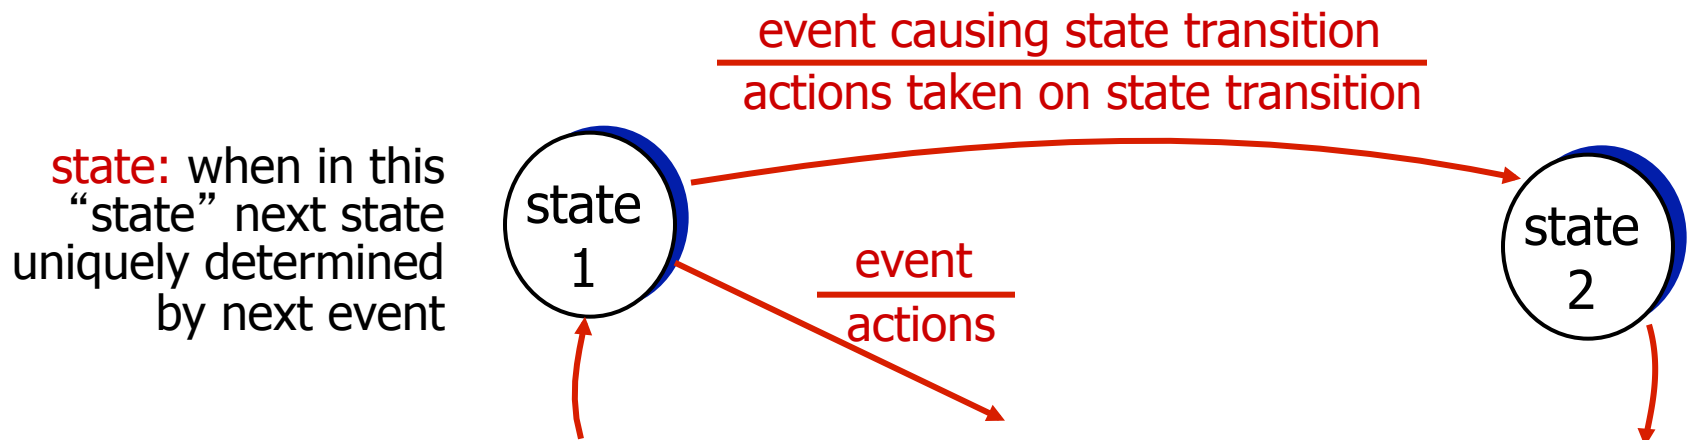
	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

Reliable data transfer: getting started

incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)

- consider only unidirectional data transfer
 - ▣ but control info will flow on both directions!
- use finite state machines (FSM) to specify sender, receiver



rdt3.0: channels with errors *and* loss

new assumption:

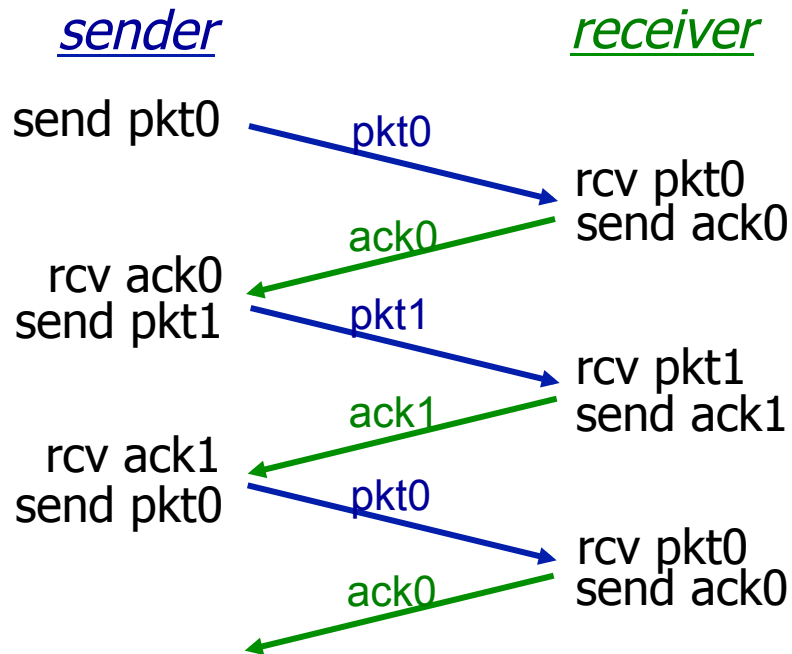
underlying channel can also lose packets (data, ACKs)

- ▣ checksum, seq. #, ACKs, retransmissions will be of help ... but not enough


approach: sender waits “reasonable” amount of time for ACK

- ▣ retransmits if no ACK received in this time
- ▣ if pkt (or ACK) just delayed (not lost):
 - ▣ retransmission will be duplicate, but seq. #'s already handles this
 - ▣ receiver must specify seq # of pkt being ACKed
- ▣ requires countdown timer

rdt3.0 in action

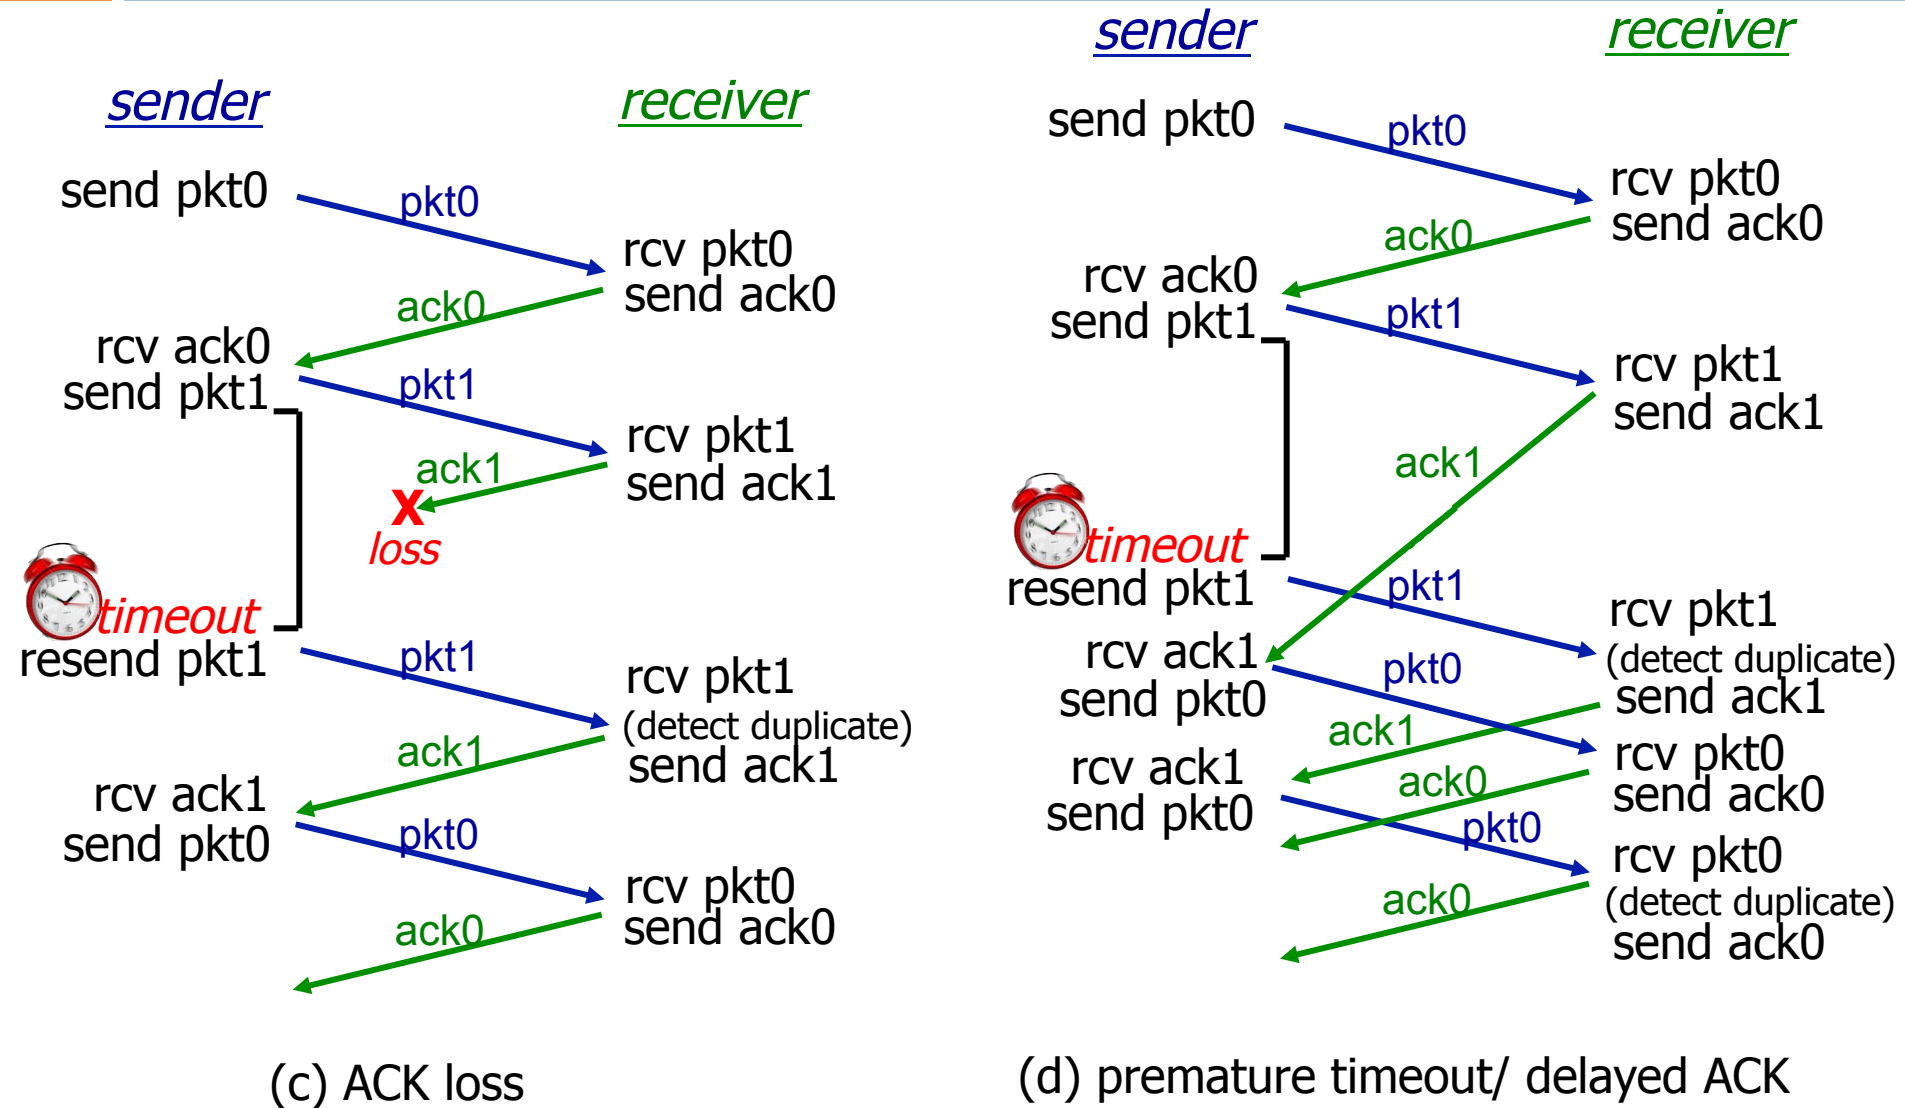


(a) no loss



(b) packet loss


rdt3.0 in action



Pipelined protocols

pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- ▣ range of sequence numbers must be increased
- ▣ buffering at sender and/or receiver

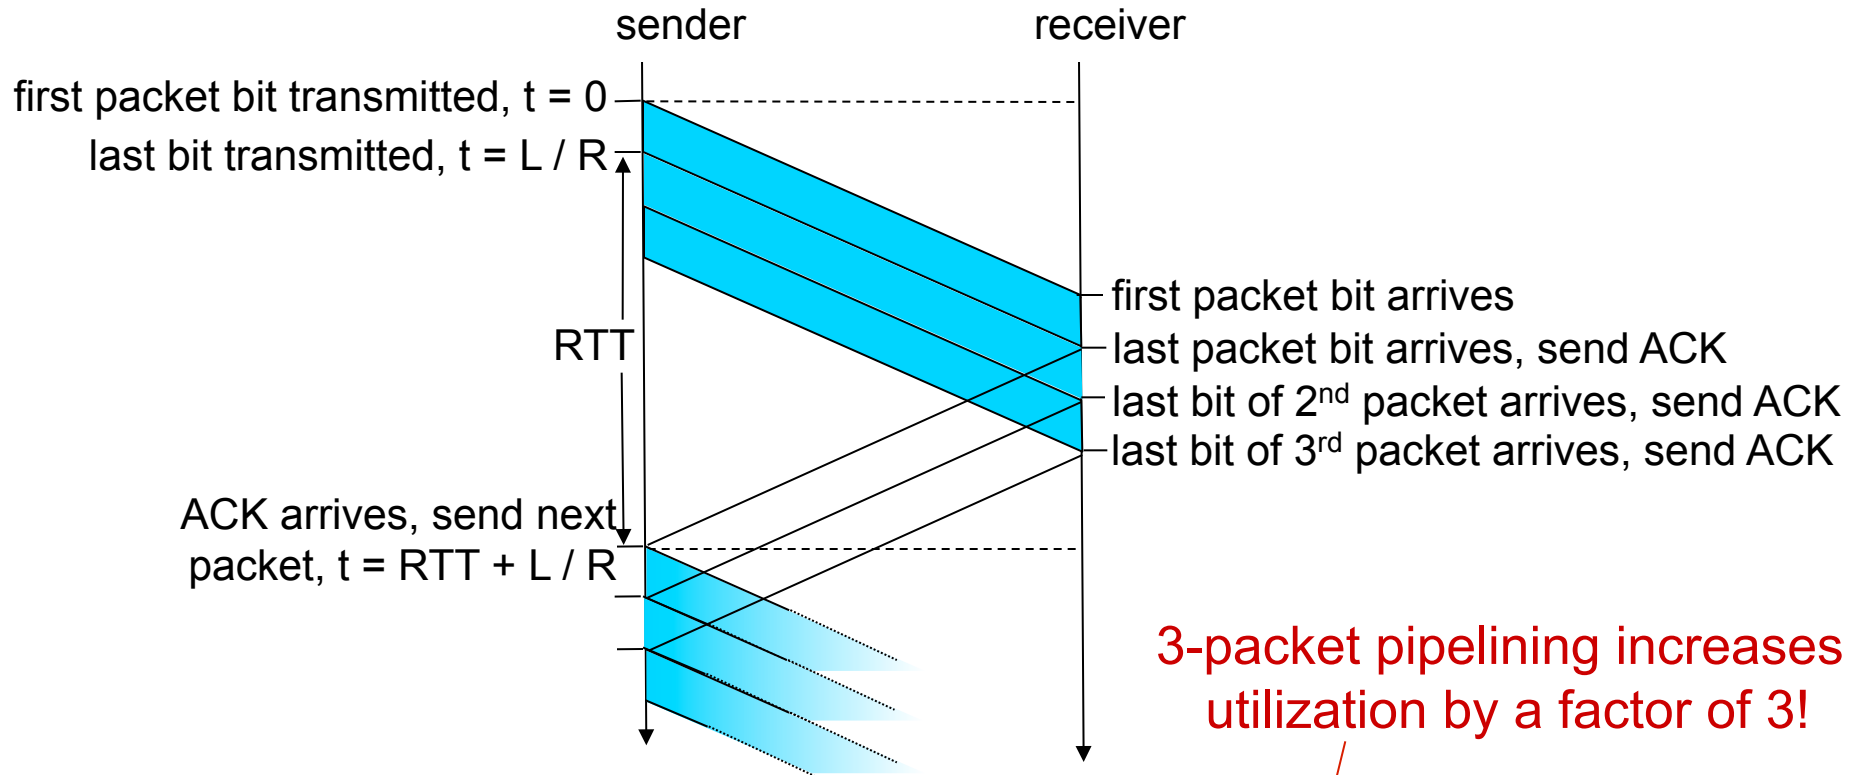


(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

- ▣ two generic forms of pipelined protocols: *go-Back-N*, *selective repeat*

Pipelining: increased utilization



3-packet pipelining increases utilization by a factor of 3!

$$U_{\text{sender}} = \frac{3L / R}{RTT + L / R} = \frac{.0024}{30.008} = 0.00081$$

Pipelined protocols: overview

Go-back-N:


- sender can have up to N unacked packets in pipeline
- receiver only sends *cumulative ack*
 - ▣ doesn't ack packet if there's a gap
- sender has timer for oldest unacked packet
 - ▣ when timer expires, retransmit *all* unacked packets

Selective Repeat:

- sender can have up to N unack'ed packets in pipeline
- rcvr sends *individual ack* for each packet
- sender maintains timer for each unacked packet
 - ▣ when timer expires, retransmit only that unacked packet

Go-Back-N: sender

- k-bit seq # in pkt header
- “window” of up to N, consecutive unack’ed pkts allowed



- ❖ ACK(n): ACKs all pkts up to, including seq # n - “*cumulative ACK*”
 - may receive duplicate ACKs (see receiver)
- ❖ timer for oldest in-flight pkt
- ❖ *timeout(n)*: retransmit packet n and all higher seq # pkts in window

Selective repeat

- receiver *individually* acknowledges all correctly received pkts
 - ▣ buffers pkts, as needed, for eventual in-order delivery to upper layer
- sender only resends pkts for which ACK not received
 - ▣ sender timer for each unACKed pkt
- sender window
 - ▣ N consecutive seq #'s
 - ▣ limits seq #'s of sent, unACKed pkts